Programming Languages Design Workshop

Vincenzo Ciancia

April 04, 2025

Contents

Chapter 1: Introduction to Programming Language Design

Section 1: What is a Programming Language?

A **programming language** is a formal language comprising a set of strings (instructions) that produce various kinds of machine output. Programming languages are used in computer programming to implement algorithms.

Key Characteristics of Programming Languages

- **Syntax**: The form or structure of the expressions, statements, and program units
- **Semantics**: The meaning of the expressions, statements, and program units
- **Type System**: The set of types and rules for how types are assigned to various constructs in the language
- **Runtime Model**: How the language executes on a computer, including memory management
- **Standard Library**: Common functionality provided out of the box

Programming Languages vs. Natural Languages

Programming languages differ from natural languages in several important ways:

1. **Precision**: Programming languages are designed to be precise and unambiguous
2. **Vocabulary**: Programming languages have a limited vocabulary defined by the language specification
3. **Grammar**: Programming languages have a strict, formal grammar with precise rules
4. **Evolution**: Programming languages evolve through explicit design decisions, not organic usage
5. **Purpose**: Programming languages are designed to instruct machines, not primarily for human communication

Section 2: The Importance of Programming Language Design

Why should we care about programming language design?

Programming Languages Shape How We Think

Programming languages are not just tools for instructing computers; they are frameworks for human thinking. Different languages emphasize different concepts and approaches:

- **Imperative languages** (C, Pascal) focus on step-by-step instructions
- **Functional languages** (Haskell, Lisp) emphasize expressions and function composition
- **Object-oriented languages** (Java, C++, Python) organize code around objects and their interactions
- **Logic languages** (Prolog) express programs as logical relations

Language Design Affects Software Quality

The design of a programming language can significantly impact:

- **Reliability**: How easy is it to write correct code?
- **Maintainability**: How easy is it to understand and modify existing code?
- **Performance**: How efficiently can the code be executed?
- **Security**: How easily can programmers avoid security vulnerabilities?
- **Developer Productivity**: How quickly can developers write and debug code?

The Evolution of Programming Languages

Programming languages have evolved dramatically over time, reflecting changes in hardware, software engineering practices, and problem domains:

- **1950s**: Assembly languages and early high-level languages (FORTRAN, LISP)
- **1960s**: ALGOL, COBOL, and structured programming concepts
- **1970s**: C, Pascal, and the rise of procedural programming
- **1980s**: C++, Ada, and the adoption of object-oriented programming

- **1990s**: Java, Python, Ruby, and the focus on portability and productivity
- **2000s**: C#, JavaScript frameworks, and web-centric languages
- **2010s**: Go, Rust, Swift, and the focus on safety and concurrency
- **2020s**: Continued evolution with AI assistance, type inference improvements, and more

Section 3: Modern Language Design Principles

What principles guide the design of modern programming languages?

Abstraction

Abstraction is the process of removing details to focus on the essential features of a concept or object.

Examples in programming languages: - Functions abstract away implementation details - Classes abstract data and behavior - Interfaces abstract expected behaviors - Modules abstract related functionality

Expressiveness

Expressiveness refers to how easily and concisely a language can express computational ideas.

Factors that contribute to expressiveness: - Rich set of operators and built-in functions - Support for higher-order functions - Pattern matching - Concise syntax for common operations

Safety

Safety features help prevent programmers from making mistakes or make it easier to find and fix errors.

Safety mechanisms in modern languages: - Static type checking - Bounds checking - Memory safety guarantees - Exception handling systems - Null safety features

Performance

Performance considerations affect how efficiently a language can be implemented and executed.

Performance factors: - Compilation vs. interpretation - Memory management approach - Static vs. dynamic typing - Optimization opportunities - Support for concurrency and parallelism

Consistency

Consistency in language design makes languages easier to learn and use correctly.

Consistency principles: - Similar concepts should have similar syntax - Minimal special cases - Orthogonal features (features that can be used in any combination) - Principle of least surprise (intuitive behavior)

Section 4: Using Python to Explore Programming Language Concepts

Why use Python for studying programming language design?

Python's Suitability for Language Implementation

Python is well-suited for implementing language interpreters and exploring language concepts:

- **Readability**: Python's clean syntax makes interpreter code easier to understand
- **High-level constructs**: Python provides lists, dictionaries, and other structures useful for language implementation
- **Dynamic typing**: Simplifies working with diverse language constructs
- **Rich standard library**: Includes parsing tools, regular expressions, and other useful utilities
- **Interactive development**: Makes experimenting with language features easier

Python 3.10+ Features Relevant to Language Design

Recent Python versions have introduced features that make it particularly interesting for PL experiments:

- **Type hints**: Allows for static type checking while maintaining dynamic execution
- **Pattern matching**: Provides elegant structural decomposition similar to functional languages
- **Dataclasses**: Simplifies creating data-carrying classes with minimal boilerplate
- **Functional programming tools**: Map, filter, reduce, lambdas, and comprehensions
- **AST module**: Allows inspection and manipulation of Python's abstract syntax tree

Section 5: Implementing Language Features in Python

Let's explore how we can implement core language components in Python.

Representing Syntax: Abstract Syntax Trees (ASTs)

An abstract syntax tree (AST) is a tree representation of the abstract syntactic structure of source code. Here's a simple example of representing expressions:

```python
from dataclasses import dataclass
from typing import Union, List

# Define the node types
@dataclass
class Number:
    value: float

@dataclass
class Variable:
    name: str

@dataclass
class BinaryOp:
    left: 'Expr'
    operator: str
    right: 'Expr'

# Define the expression type
Expr = Union[Number, Variable, BinaryOp]

# Example: 2 + (x * 3)
expr = BinaryOp(
    left=Number(2.0),
    operator='+',
    right=BinaryOp(
        left=Variable('x'),
        operator='*',
        right=Number(3.0)
    )
)
```

Implementing an Evaluator

The evaluator traverses the AST and computes the result. For example:

```python
def evaluate(expr: Expr, environment: dict = None) -> float:
    """Evaluate an expression in the given environment."""
    if environment is None:
        environment = {}

    if isinstance(expr, Number):
        return expr.value
    elif isinstance(expr, Variable):
        if expr.name not in environment:
            raise NameError(f"Variable '{expr.name}' not defined")
        return environment[expr.name]
    elif isinstance(expr, BinaryOp):
        left_val = evaluate(expr.left, environment)
        right_val = evaluate(expr.right, environment)

        if expr.operator == '+':
            return left_val + right_val
        elif expr.operator == '-':
            return left_val - right_val
        elif expr.operator == '*':
            return left_val * right_val
        elif expr.operator == '/':
            return left_val / right_val
        else:
            raise ValueError(f"Unknown operator: {expr.operator}")
```

Pattern Matching for AST Processing

Python 3.10's pattern matching provides a more elegant way to implement evaluators:

```python
def evaluate_with_match(expr: Expr, environment: dict = None) -> float:
    """Evaluate an expression using pattern matching."""
    if environment is None:
        environment = {}

    match expr:
        case Number(value):
            return value
        case Variable(name):
            if name not in environment:
                raise NameError(f"Variable '{name}' not defined")
            return environment[name]
        case BinaryOp(left, operator, right):
            left_val = evaluate_with_match(left, environment)
            right_val = evaluate_with_match(right, environment)

            match operator:
                case '+':
                    return left_val + right_val
                case '-':
                    return left_val - right_val
                case '*':
                    return left_val * right_val
                case '/':
```

```
                return left_val / right_val
            case _:
                raise ValueError(f"Unknown operator: {operator}")
```

Simple Type Checking

We can implement basic type checking for our language:

```python
from enum import Enum, auto
from dataclasses import dataclass
from typing import Dict

class Type(Enum):
    NUMBER = auto()
    BOOLEAN = auto()
    STRING = auto()

def type_check(expr: Expr, type_env: Dict[str, Type]) -> Type:
    """Determine the type of an expression."""
    match expr:
        case Number(_):
            return Type.NUMBER
        case Variable(name):
            if name not in type_env:
                raise TypeError(f"Variable '{name}' not defined")
            return type_env[name]
        case BinaryOp(left, operator, right):
            left_type = type_check(left, type_env)
            right_type = type_check(right, type_env)

            # Type checking rules for operators
            if operator in ['+', '-', '*', '/']:
                if left_type != Type.NUMBER or right_type != Type.NUMBER:
                    raise TypeError(f"Operator '{operator}' requires number operands")
                return Type.NUMBER
            elif operator in ['==', '!=', '<', '>', '<=', '>=']:
                if left_type != right_type:
                    raise TypeError("Comparison operators require operands of the same type")
                return Type.BOOLEAN
            else:
                raise ValueError(f"Unknown operator: {operator}")
```

Section 6: Course Structure

This course will introduce you to programming language design concepts through hands-on implementation in Python.

Course Topics

Throughout this course, we will cover:

1. **Language Syntax and Semantics**
   - Parsing and lexical analysis
   - Abstract syntax trees
   - Operational semantics
2. **Type Systems**

- Static vs. dynamic typing
- Type inference
- Polymorphism
- Advanced type features (generics, algebraic data types)
3. **Language Features**
   - Functions and closures
   - Pattern matching
   - Object-oriented programming
   - Concurrency models
   - Memory management approaches
4. **Interpreter and Compiler Implementation**
   - Building a simple interpreter
   - Environment and scope
   - Evaluation strategies
   - Introduction to compilation concepts

Projects and Exercises

The course will include:

- Regular programming exercises to reinforce concepts
- Progressive development of a language interpreter
- Exploration of existing language implementations
- Analysis of language design trade-offs

Section 7: Prerequisites and Setup

Knowledge Prerequisites

To get the most out of this course, you should have:

- Basic Python programming experience
- Understanding of fundamental programming concepts (variables, functions, control flow)
- Familiarity with basic data structures (lists, dictionaries, trees)
- Interest in how programming languages work "under the hood"

No prior experience with compiler or interpreter development is required.

Python Environment Setup

To follow along with the course examples and exercises:

1. **Install Python 3.10 or later**
   - Required for pattern matching and other modern features
2. **Recommended development tools**
   - Visual Studio Code with Python extension
   - PyCharm
   - Jupyter Notebook/Lab for interactive exploration
3. **Useful libraries**
   - mypy for static type checking
   - pytest for testing your implementations

Section 8: Additional Resources

Books on Programming Language Design

- **"Crafting Interpreters"** by Robert Nystrom

- **"Programming Language Pragmatics"** by Michael Scott
- **"Types and Programming Languages"** by Benjamin Pierce
- **"Concepts of Programming Languages"** by Robert Sebesta
- **"Structure and Interpretation of Computer Programs"** by Abelson and Sussman

Online Resources

- Python Documentation
- Python Type Hints
- Pattern Matching in Python 3.10
- The AST Module
- Building a Simple Interpreter

Chapter 2: Types and Pattern Matching in Python

Section 1: Python's Type System

What are Type Annotations?

Python's type system allows developers to add optional type hints to variables, function parameters, and return values. These annotations help catch errors early, improve code documentation, and enhance IDE support without changing the runtime behavior of the code.

```python
def greet(name: str) -> str:
    return f"Hello, {name}!"
```

Type annotations are part of the Python Enhancement Proposal (PEP) 484 and have been continuously improved in subsequent PEPs. They provide a way to make Python code more robust through static type checking, although Python remains dynamically typed at runtime.

Basic Types in Python

Python provides several built-in types for annotations:

- **Primitive types**: int, float, bool, str
- **Collection types**: list, tuple, dict, set

Example from types_and_matching.py:

```python
my_tuple: tuple[int, str, float] = (1, "hello", 1.0)
my_list: list[int] = [1, 2, 3, 4, 5]
```

These annotations tell the type checker that my_tuple is a 3-element tuple containing an integer, a string, and a float, while my_list is a list containing only integers.

Generic Types

Generic types allow you to create reusable, type-safe components. In Python 3.12+, the syntax for generic classes uses square brackets:

```python
class Stack[T]:
    def __init__(self) -> None:
        self.items: list[T] = []

    def push(self, item: T) -> None:
        self.items.append(item)

    def pop(self) -> T | None:
        return self.items.pop() if self.items else None
```

In this example from our code: - T is a type parameter representing any type - Stack[T] is a generic class that can be specialized for specific types - list[T] indicates a list containing elements of type T

To use this generic class:

```python
stack_1 = Stack[int]()  # A stack of integers
stack_1.push(1)  # Valid
stack_1.push("hello")  # Type error: expected int, got str
```

Union Types and Optional Values

Union types allow a variable to have multiple possible types, expressed using the | operator (introduced in Python 3.10):

```python
def process_data(data: int | str) -> None:
    # Can handle both integers and strings
    pass
```

From our example code:

```python
def pop(self) -> T | None:  # Use | for union types
    return self.items.pop() if self.items else None
```

This indicates that the `pop` method returns either a value of type `T` or `None` if the stack is empty.

Type Aliases

Type aliases help simplify complex type annotations:

```python
type JsonData = dict[str, int | str | list | dict]
type A = int | dict[str, A]  # Recursive type
```

In `types_and_matching.py`, we use type aliases for recursive types:

```python
type MyBaseList[T] = None | MyList[T]
type Expr = int | Sum
```

These aliases make the code more readable and allow for recursive type definitions.

Literal Types

Literal types restrict values to specific constants:

```python
y: Literal["hello", "world"] = "hello"  # Valid
z: Literal[1, 2, 3] = 9  # Type error
```

From our example code:

```python
@dataclass
class Human:
    name: str
    drivingLicense: Literal[True, False]
```

This constrains the `drivingLicense` attribute to be either `True` or `False` only.

Section 2: Structural Pattern Matching

Introduction to Pattern Matching

Introduced in Python 3.10, the `match` statement provides powerful pattern matching capabilities, similar to switch statements in other languages but with more expressive power:

```python
def describe(value):
    match value:
        case 0:
            return "Zero"
        case int(x) if x > 0:
            return "Positive integer"
        case str():
            return "String"
        case _:
            return "Something else"
```

Pattern matching allows for more concise and readable code, especially when dealing with complex data structures and multiple conditions.

Basic Patterns

Basic patterns match against simple values and types:

```python
match x:
    case 0:
        print("Zero")
    case int():
        print("Integer")
    case str():
        print("String")
    case _:  # Wildcard pattern
        print("Something else")
```

The _ pattern is a wildcard that matches anything and is often used as a catch-all case.

Sequence Patterns

Sequence patterns match against sequence types like lists and tuples:

```python
def process_lst(lst: list[int]) -> None:
    match lst:
        case []:
            print("List: empty")
        case [head, *tail]:
            print(f"List: head {head}, tail {tail}")
```

This example shows destructuring a list into its head (first element) and tail (remaining elements), demonstrating how pattern matching facilitates recursive list processing.

Class Patterns and Attribute Matching

Pattern matching works particularly well with dataclasses:

```python
def greet(person: Person | str) -> None:
    match person:
        case Person(name="Alice", age=25):
            print("Hello Alice!")
        case Person(name=x, age=y):
            print(f"Hello {x} of age {y}!")
        case str():
            print(f"Hello {person}!")
```

This example shows matching against specific attribute values and binding attributes to variables.

Complex Pattern Matching Examples

Recursive Pattern Matching   Pattern matching excels at handling recursive data structures:

```python
def sum_list_2(lst: MyBaseList[int]) -> int:
    match lst:
        case None:
            return 0
        case MyList(head=x, tail=y):
            return x + sum_list_2(y)
```

This function processes a custom linked list structure using pattern matching to handle the base case (None) and recursive case elegantly.

Parsing Expressions   Pattern matching can implement simple interpreters:

```python
def eval_expr(expr: Expr) -> int:
    match expr:
        case int(x):
            return x
        case Sum(left=x, right=y):
            return eval_expr(x) + eval_expr(y)
```

This code evaluates a simple arithmetic expression tree, showing how pattern matching simplifies traversal of complex structures.

Section 3: Combining Types and Pattern Matching

Algebraic Data Types in Python

Python can implement algebraic data types (ADTs) using classes, dataclasses, and union types:

Sum Types (Tagged Unions)   Sum types represent values that could be one of several alternatives:

```python
@dataclass
class Human:
    name: str
    drivingLicense: Literal[True, False]


@dataclass
class Dog:
    name: str
    kind: DogKind
    colour: Literal["brown", "black", "white"]


type Record = Human | Dog
```

This example defines two distinct record types (Human and Dog) and a union type Record that can be either of them.

Pattern Matching with Sum Types   Pattern matching works seamlessly with sum types:

```python
def describe_person(person: Human | Dog) -> str:
    match person:
        case Human(name=name, drivingLicense=True):
            return f"Person {name} has a driving license"
        case Human(name=name, drivingLicense=False):
            return f"Person {name} does not have a driving license"
        case Dog(name=name, kind=kind, colour=colour):
            return f"Dog {name} is a {kind} and has a {colour} colour"
```

This function handles different record types with specific patterns for each case.

Product Types   Product types represent combinations of values:

```python
@dataclass
class Person:
    name: str
    age: int
```

A dataclass like Person is a product type, representing a combination of a string and an integer.

Type Checking and Pattern Matching

Type checkers like mypy can catch errors in pattern matching code:

```python
def describe_person_2(person: Record) -> str:
    if isinstance(person, Human):
        return person.name + person.colour  # Type error: Human has no attribute 'colour'
    else:
        return person.name + person.kind
```

Type checking helps identify incorrect attribute access, which might otherwise lead to runtime errors.

Section 4: Applications and Best Practices

When to Use Type Annotations

Type annotations are particularly valuable in: - Large codebases with multiple developers - APIs and libraries that will be used by others - Performance-critical code where type-specific optimizations matter - Complex data processing pipelines

Domain-specific languages, interpreters, ADTs!!

When to Use Pattern Matching

Pattern matching excels at: - Processing complex recursive data structures - Implementing interpreters and compilers - Handling case-based logic with destructuring - Processing structured data like JSON or ASTs

Best Practices for Type Annotations

1. Be consistent with type annotations across your codebase
2. Use type aliases for complex or repetitive type expressions
3. Leverage tools like mypy, pyright, or pylance for static type checking
4. Balance between type precision and code readability
5. Document non-obvious type constraints with comments

Best Practices for Pattern Matching

1. Order cases from most specific to most general
2. Use the wildcard pattern (_) as the last case
3. Consider using guard clauses for complex conditions
4. Break complex pattern matching into smaller functions
5. Leverage destructuring to avoid redundant variable assignments

Exercises

- Consider the expression evaluator written in Lecture 01. Turn it into a parser that converts a string of the form "x op y op z …" separated by spaces into an AST in the sense of Lecture 02. Concatenate the new parser and the evaluation function that takes AST as input to define a mini-interpreter.

- Add the "==" boolean operator to the AST and the evaluator.

Additional Resources

- PEP 484 – Type Hints
- PEP 585 – Type Hinting Generics In Standard Collections
- PEP 604 – Allow writing union types as X | Y
- PEP 634 – Structural Pattern Matching: Specification
- PEP 636 – Structural Pattern Matching: Tutorial

- Mypy Type Checker Documentation
- Real Python: Python Type Checking
- Real Python: Structural Pattern Matching in Python

Chapter 3: Building a Mini Interpreter

Section 1: Introduction to Interpreters

An interpreter is a program that executes source code directly, without requiring compilation to machine code. Let's explore how interpreters work and how to build a simple one in Python.

The Role of Interpreters

Interpreters serve several important purposes in programming language implementation:

1. **Direct Execution**: Execute source code without a separate compilation step
2. **Immediate Feedback**: Provide instant results for interactive programming
3. **Portability**: Run on any platform that supports the interpreter
4. **Simplicity**: Often easier to implement than full compilers
5. **Debugging**: Allow for interactive debugging and inspection

Languages like Python, JavaScript, and Ruby primarily use interpreters, while others like Java use a hybrid approach with compilation to bytecode followed by interpretation.

Components of an Interpreter

A typical interpreter includes the following components:

1. **Lexer (Tokenizer)**: Converts source code text into tokens
2. **Parser**: Transforms tokens into an abstract syntax tree (AST)
3. **Evaluator**: Executes the AST to produce results
4. **Environment**: Stores variables and their values
5. **Error Handler**: Manages and reports errors

We'll implement each of these components in our mini interpreter.

Section 2: Lexical Analysis

The first step in interpreting code is breaking it down into tokens - the smallest meaningful units in the language.

Tokens and Lexical Structure

Tokens are the building blocks of a language, similar to words in natural language. Common token types include:

- **Keywords**: Reserved words with special meaning (e.g., `if`, `while`)
- **Identifiers**: Names given to variables, functions, etc.
- **Literals**: Constant values (numbers, strings, booleans)
- **Operators**: Symbols that perform operations (`+`, `-`, `*`, `/`)
- **Punctuation**: Symbols that structure the code (`;`, `,`, `{}`, `()`)

Implementing a Lexer

Our lexer will convert a string of source code into a list of tokens:

```python
from dataclasses import dataclass
from enum import Enum, auto
from typing import List, Optional


# Define token types
class TokenType(Enum):
    NUMBER = auto()
    PLUS = auto()
```

```python
    MINUS = auto()
    MULTIPLY = auto()
    DIVIDE = auto()
    LPAREN = auto()
    RPAREN = auto()
    EOF = auto()  # End of file

@dataclass
class Token:
    type: TokenType
    value: Optional[str] = None

def tokenize(text: str) -> List[Token]:
    tokens = []
    i = 0

    while i < len(text):
        char = text[i]

        # Skip whitespace
        if char.isspace():
            i += 1
            continue

        # Process numbers
        if char.isdigit():
            num = ""
            while i < len(text) and text[i].isdigit():
                num += text[i]
                i += 1
            tokens.append(Token(TokenType.NUMBER, num))
            continue

        # Process operators
        if char == '+':
            tokens.append(Token(TokenType.PLUS))
        elif char == '-':
            tokens.append(Token(TokenType.MINUS))
        elif char == '*':
            tokens.append(Token(TokenType.MULTIPLY))
        elif char == '/':
            tokens.append(Token(TokenType.DIVIDE))
        elif char == '(':
            tokens.append(Token(TokenType.LPAREN))
        elif char == ')':
            tokens.append(Token(TokenType.RPAREN))
        else:
            raise ValueError(f"Unexpected character: {char}")

        i += 1

    tokens.append(Token(TokenType.EOF))
    return tokens
```

Testing the Lexer

Let's test our lexer with a simple arithmetic expression:

```python
def test_lexer():
    source = "3 + 4 * (2 - 1)"
    tokens = tokenize(source)

    expected = [
        Token(TokenType.NUMBER, "3"),
        Token(TokenType.PLUS),
        Token(TokenType.NUMBER, "4"),
        Token(TokenType.MULTIPLY),
        Token(TokenType.LPAREN),
        Token(TokenType.NUMBER, "2"),
        Token(TokenType.MINUS),
        Token(TokenType.NUMBER, "1"),
        Token(TokenType.RPAREN),
        Token(TokenType.EOF)
    ]

    assert tokens == expected, f"Expected {expected}, got {tokens}"
    print("Lexer test passed!")

test_lexer()
```

Section 3: Parsing and Abstract Syntax Trees

Once we have tokens, we need to organize them into a structured representation of the program - an Abstract Syntax Tree (AST).

Understanding Abstract Syntax Trees

An AST is a tree representation of the abstract syntactic structure of source code. Each node in the tree represents a construct in the source code.

For example, the expression `3 + 4 * 2` would be represented as:

```
    (+)
   /   \
  3    (*)
      /   \
     4     2
```

The tree captures the structure and precedence of operations.

Defining AST Nodes

Let's define classes for our AST nodes:

```python
from dataclasses import dataclass
from typing import Union, Optional


@dataclass
class Number:
    value: float


@dataclass
```

19

```python
class BinaryOp:
    left: 'Expression'
    operator: str
    right: 'Expression'


# Define our Expression type
Expression = Union[Number, BinaryOp]
```

Implementing a Recursive Descent Parser

A recursive descent parser is a top-down parser that uses a set of recursive procedures to process the input:

```python
class Parser:
    def __init__(self, tokens: List[Token]):
        self.tokens = tokens
        self.current = 0

    def peek(self) -> Token:
        return self.tokens[self.current]

    def previous(self) -> Token:
        return self.tokens[self.current - 1]

    def advance(self) -> Token:
        if not self.is_at_end():
            self.current += 1
        return self.previous()

    def is_at_end(self) -> bool:
        return self.peek().type == TokenType.EOF

    def check(self, type: TokenType) -> bool:
        if self.is_at_end():
            return False
        return self.peek().type == type

    def match(self, *types: TokenType) -> bool:
        for type in types:
            if self.check(type):
                self.advance()
                return True
        return False

    def consume(self, type: TokenType, message: str) -> Token:
        if self.check(type):
            return self.advance()
        raise Exception(message)

    def parse(self) -> Expression:
        return self.expression()

    def expression(self) -> Expression:
        return self.term()
```

```python
    def term(self) -> Expression:
        expr = self.factor()

        while self.match(TokenType.PLUS, TokenType.MINUS):
            operator = self.previous().type
            right = self.factor()
            expr = BinaryOp(
                expr,
                "+" if operator == TokenType.PLUS else "-",
                right
            )

        return expr

    def factor(self) -> Expression:
        expr = self.primary()

        while self.match(TokenType.MULTIPLY, TokenType.DIVIDE):
            operator = self.previous().type
            right = self.primary()
            expr = BinaryOp(
                expr,
                "*" if operator == TokenType.MULTIPLY else "/",
                right
            )

        return expr

    def primary(self) -> Expression:
        if self.match(TokenType.NUMBER):
            return Number(float(self.previous().value))

        if self.match(TokenType.LPAREN):
            expr = self.expression()
            self.consume(TokenType.RPAREN, "Expect ')' after expression.")
            return expr

        raise Exception(f"Unexpected token: {self.peek()}")
```

This parser implements the following grammar:

```
expression → term
term       → factor (("+"|"-") factor)*
factor     → primary (("*"|"/") primary)*
primary    → NUMBER | "(" expression ")"
```

Testing the Parser

Let's test our parser with the same expression:

```python
def test_parser():
    tokens = tokenize("3 + 4 * 2")
    parser = Parser(tokens)
    ast = parser.parse()
```

```
    # Expected: BinaryOp(Number(3), "+", BinaryOp(Number(4), "*", Number(2)))
    expected = BinaryOp(
        Number(3),
        "+",
        BinaryOp(
            Number(4),
            "*",
            Number(2)
        )
    )

    assert ast == expected, f"Expected {expected}, got {ast}"
    print("Parser test passed!")


test_parser()
```

Section 4: Evaluating Expressions

Now that we have an AST, we can evaluate it to produce a result.

The Evaluation Process

Evaluation is the process of computing the result of an expression. It typically involves:

1. Walking the AST recursively
2. Computing the value of each node based on its type and children
3. Combining results according to the language semantics

Implementing an Evaluator

For our mini interpreter, we'll implement a simple evaluator that computes arithmetic expressions:

```python
def evaluate(expr: Expression) -> float:
    """Evaluate an expression and return its value."""
    match expr:
        case Number(value):
            return value
        case BinaryOp(left, operator, right):
            left_value = evaluate(left)
            right_value = evaluate(right)

            match operator:
                case "+":
                    return left_value + right_value
                case "-":
                    return left_value - right_value
                case "*":
                    return left_value * right_value
                case "/":
                    if right_value == 0:
                        raise ZeroDivisionError("Division by zero")
                    return left_value / right_value
                case _:
                    raise ValueError(f"Unknown operator: {operator}")
```

Testing the Evaluator

Let's test our evaluator with a few expressions:

```python
def test_evaluator():
    expressions = [
        ("3 + 4", 7),
        ("3 * 4", 12),
        ("10 - 2", 8),
        ("20 / 5", 4),
        ("3 + 4 * 2", 11),   # Tests operator precedence
        ("(3 + 4) * 2", 14),   # Tests parentheses
    ]

    for source, expected in expressions:
        tokens = tokenize(source)
        parser = Parser(tokens)
        ast = parser.parse()
        result = evaluate(ast)

        assert result == expected, f"For '{source}', expected {expected}, got {result}"

    print("Evaluator tests passed!")

test_evaluator()
```

Section 5: Putting It All Together

Now we'll combine our lexer, parser, and evaluator into a complete mini interpreter.

The Full Interpreter

```python
def interpret(source: str) -> float:
    """Interpret a source string and return the result."""
    try:
        tokens = tokenize(source)
        parser = Parser(tokens)
        ast = parser.parse()
        return evaluate(ast)
    except Exception as e:
        print(f"Error: {e}")
        return float('nan')  # Return NaN for errors

def run_repl():
    """Run a simple Read-Eval-Print Loop."""
    print("Mini Interpreter REPL")
    print("Enter expressions to evaluate, or 'exit' to quit")

    while True:
        try:
            source = input("> ")
            if source.lower() in ("exit", "quit"):
                break

            result = interpret(source)
```

```
        print(f"= {result}")
    except KeyboardInterrupt:
        break
    except Exception as e:
        print(f"Error: {e}")


    print("Goodbye!")


if __name__ == "__main__":
    run_repl()
```

Example Usage

Using our interpreter:

```
Mini Interpreter REPL
Enter expressions to evaluate, or 'exit' to quit
> 3 + 4
= 7.0
> 3 * (4 + 2)
= 18.0
> 10 / (2 - 2)
Error: Division by zero
> (3 + 4) * (5 - 2)
= 21.0
> exit
Goodbye!
```

Section 6: Extending the Interpreter

Our mini interpreter is very basic, but we can extend it with more features.

Adding Variables

To add variable support, we need:

1. An environment to store variable bindings
2. New AST nodes for variable references and assignments
3. Updated parsing rules
4. Updated evaluation logic

Let's implement a simple environment first:

```
@dataclass
class Environment:
    values: dict[str, float] = None

    def __post_init__(self):
        if self.values is None:
            self.values = {}

    def define(self, name: str, value: float) -> None:
        """Define a variable with the given name and value."""
        self.values[name] = value

    def get(self, name: str) -> float:
```

```python
        """Get the value of a variable."""
        if name in self.values:
            return self.values[name]
        raise ValueError(f"Undefined variable: {name}")
```

New AST Nodes for Variables

```python
@dataclass
class Variable:
    name: str


@dataclass
class Assignment:
    name: str
    value: 'Expression'


# Update Expression type
Expression = Union[Number, BinaryOp, Variable, Assignment]
```

Updating the Parser

```python
class Parser:
    # ... existing code ...

    def parse(self) -> Expression:
        return self.assignment()

    def assignment(self) -> Expression:
        expr = self.expression()

        if self.match(TokenType.EQUAL):
            if isinstance(expr, Variable):
                value = self.assignment()
                return Assignment(expr.name, value)
            raise Exception("Invalid assignment target")

        return expr

    # ... update tokenize() to handle identifiers and '=' ...
```

Updating the Evaluator

```python
def evaluate(expr: Expression, env: Environment) -> float:
    """Evaluate an expression in the given environment."""
    match expr:
        case Number(value):
            return value
        case Variable(name):
            return env.get(name)
        case Assignment(name, value):
            result = evaluate(value, env)
            env.define(name, result)
            return result
```

```
        case BinaryOp(left, operator, right):
            # ... existing code ...
```

Adding Control Flow

We can extend our language with if-expressions:

```
@dataclass
class If:
    condition: Expression
    then_branch: Expression
    else_branch: Expression


# Update Expression type
Expression = Union[Number, BinaryOp, Variable, Assignment, If]


# Update evaluator
def evaluate(expr: Expression, env: Environment) -> float:
    match expr:
        # ... existing cases ...
        case If(condition, then_branch, else_branch):
            if evaluate(condition, env) != 0:  # Non-zero is true
                return evaluate(then_branch, env)
            else:
                return evaluate(else_branch, env)
```

Section 7: Further Exploration

Here are some ways you could extend our mini interpreter further:

Additional Features to Implement

1. **Functions**: Add function declarations and calls
2. **Loops**: Implement while or for loops
3. **More Operators**: Add comparison, logical, and bitwise operators
4. **Error Handling**: Improve error messages and recovery
5. **Type System**: Add a simple type system
6. **Standard Library**: Implement built-in functions for common operations

Learning Resources

To learn more about interpreters and language implementation:

- **"Crafting Interpreters"** by Robert Nystrom: A comprehensive guide to implementing interpreters
- **"Programming Language Pragmatics"** by Michael Scott: Covers theoretical aspects of language design
- **"Structure and Interpretation of Computer Programs"** by Abelson and Sussman: A classic text on programming language concepts

Conclusion

Building a mini interpreter helps understand how programming languages work under the hood. We've seen how to:

1. Tokenize source code into lexical tokens
2. Parse tokens into an abstract syntax tree
3. Evaluate the AST to produce results
4. Extend the interpreter with new features

This foundation can be expanded to build more complex languages and tools.

Chapter 4: Semantic Domains and Environment-Based Interpreters

Section 1: Introduction to Semantic Domains

In programming language semantics, **semantic domains** are mathematical structures used to give meaning to syntactic constructs. They provide the foundation for defining the behavior of programs in a precise, mathematical way.

Key Semantic Domains in Programming Languages

- **Syntactic Domains**: Represent the structure of programs (tokens, parse trees, syntax trees).
- **Semantic Domains**: Represent the meaning of programs (values, environments, states).

The relationship between these domains is at the heart of language semantics:

```
Syntax -> Semantics
Program -> Meaning
```

Core Semantic Domains

In our mini-interpreter, we'll work with several fundamental semantic domains:

- **Expressible Values**: Values that can be produced by evaluating expressions
- **Denotable Values (DVal)**: Values that can be bound to identifiers in the environment
- **Memorizable Values (MVal)**: Values that can be stored in memory/state
- **Identifiers**: Names of constants, variables, functions, modules, etc.
- **Locations**: Memory addresses of variables
- **Environment**: Maps identifiers to denotable values, representing the binding context
- **State**: Maps memory locations to memorizable values, representing the program's memory

While these domains often overlap, they aren't necessarily identical. Understanding the differences is crucial for language design.

Side Effects and Pure Functions

A fundamental distinction in programming language semantics is between **pure** computations and those that produce **side effects**.

Pure Functions   A **pure function** is a computation that: 1. Always produces the same output for the same input 2. Has no observable effects beyond computing its result

In mathematical terms, a pure function is simply a mapping from inputs to outputs, like mathematical functions (e.g., $\sin(x)$, $\log(x)$).

```python
def add(x: int, y: int) -> int:
    return x + y  # Pure: same inputs always give same output
```

Side Effects   A **side effect** is any observable change to the system state that occurs during computation, beyond returning a value. Common side effects include:

1. **Memory updates**: Modifying variables or data structures

   ```python
   x = 10  # Changes the program's state
   ```

   Q: in this context, what is a function that does **not** return always the same value for the same inputs?

2. **Input/Output operations**:

   ```python
   print("Hello")  # Affects the external world (terminal)
   ```

3. **File operations**:

```python
with open("data.txt", "w") as f:
    f.write("data")  # Changes the file system
```

4. **Network operations**:

```python
requests.get("https://example.com")  # Interacts with external systems
```

Memory Updates as Side Effects   In our interpreter design, memory (state) updates are a primary form of side effect. When we update state:

```python
def state_update(state: State, location: int, value: MVal) -> State:
    new_state = state.copy()
    new_state[location] = value
    return new_state
```

We're representing a change to the program's memory. In a real computer, this would modify memory cells directly. Our functional implementation returns a new state rather than modifying the existing one, but conceptually it represents the same side effect.

Side Effects in Programming Languages   Languages differ in how they handle side effects:

- **Purely functional languages** (like Haskell) isolate side effects using type systems and monads
- **Imperative languages** (like C, Python) embrace side effects as their primary mechanism for computation
- **Hybrid languages** (like Scala, OCaml) support both styles

Understanding side effects is crucial for language design because they impact: - Program correctness (pure functions are easier to reason about) - Parallelization (side effects complicate parallel execution) - Optimization (pure functions allow more aggressive optimizations)

In our interpreter, we'll model side effects using explicit state passing, maintaining the mathematical clarity of our semantics while accurately representing the behavior of stateful programs.

Section 2: Denotable vs. Memorizable Values

Denotable Values (DVal)

Denotable values are those that can be bound to identifiers in an environment. In our implementation, they include:

```python
type Num = int  # A type alias for integers
type DenOperator = Callable[[int, int], int]
type DVal = int | DenOperator  # Denotable values
```

DVal includes: - **Numbers**: Simple integer values - **Operators**: Functions that take two integers and return an integer

Memorizable Values (MVal)

Memorizable values are those that can be stored in memory (the state). In our implementation:

```python
type MVal = int  # Memorizable values
```

MVal only includes integers, not functions. This highlights an important distinction:

> Not everything that can be bound to a name can be stored in memory.

This distinction is crucial for understanding: - Why some languages don't support first-class functions - Why some types require special treatment in memory management - How languages with different type systems handle values differently

Section 3: Environment and State as Functions

In our treatment of semantic domains, we adopt a purely functional approach where environments are represented as functions, and states are represented as dataclasses containing store functions, rather than mutable data structures. This approach aligns with the mathematical view of semantic domains and provides a clean conceptual model for understanding program behavior.

Functional Programming: A Brief Digression

Before diving into our function-based implementation of environments and state, it's worth taking a brief detour to discuss functional programming concepts, as they form the foundation of our approach.

Functional programming is a paradigm where computations are treated as evaluations of mathematical functions, emphasizing immutable data and avoiding side effects. Python, while not a pure functional language, supports many functional programming techniques.

Functions as First-Class Citizens   In functional programming, functions are "first-class citizens" — they can be: - Assigned to variables - Passed as arguments to other functions - Returned from functions - Stored in data structures

For example:

```python
# Function assigned to a variable
increment = lambda x: x + 1


# Function passed as an argument
def apply_twice(f, x):
    return f(f(x))


result = apply_twice(increment, 3)  # Returns 5
```

Higher-Order Functions: Map   A common pattern in functional programming is applying a function to each element in a collection. Python's `map` function does exactly this:

```python
numbers = [1, 2, 3, 4, 5]


# Apply a function to each element
squared = list(map(lambda x: x**2, numbers))  # [1, 4, 9, 16, 25]


# Equivalent to a list comprehension
squared_alt = [x**2 for x in numbers]  # [1, 4, 9, 16, 25]
```

Function Composition   Functional programming emphasizes building complex behaviors by composing simpler functions:

```python
def compose(f, g):
    return lambda x: f(g(x))


# Compose two functions
negate_and_square = compose(lambda x: -x, lambda x: x**2)
result = negate_and_square(5)  # -25
```

Pure Functions and Immutability   Pure functions always produce the same output for the same input and have no side effects. This property makes them predictable and easier to reason about:

```python
# Pure function
def add(a, b):
    return a + b
```

```
# Impure function (has side effects)
def add_and_print(a, b):
    result = a + b
    print(f"The result is {result}")  # Side effect: printing
    return result
```

Relevance to Semantic Domains   These functional programming concepts directly inform our approach to implementing semantic domains:

1. We represent environments and stores as functions, not data structures
2. We use higher-order functions to create updated environments and stores
3. We maintain immutability through functional updates rather than mutations
4. We compose simple operations to build complex behaviors

With this foundation in mind, let's explore how we represent environments and states as functions.

Environment as a Function

An environment is mathematically a function that maps identifiers to denotable values:

```
type Environment = Callable[[str], DVal]
```

This means an environment is a function that: - Takes an identifier (string) as input - Returns a denotable value (DVal) - Raises an error if the identifier is not defined

While many practical implementations use dictionaries or hash tables for efficiency, conceptually an environment is simply a function:

```
Environment: Identifier → DVal
```

State as a Dataclass

Unlike environment, which is purely a function, a state encapsulates both a store function and allocation information:

```
@dataclass
class State:
    store: Callable[[Location], MVal]  # The store function
    next_loc: int                      # Next available location
```

This represents a state as: - A dataclass containing a store function and next location counter - The store function takes a location as input and returns the value at that location - The next_loc field tracks the next available memory location for allocation

Conceptually, the store component maintains the mapping:

```
Store: Location → MVal
```

While the next_loc component tracks allocation state.

Section 4: Functional Updates

In a purely functional approach, we don't mutate existing environments or states. Instead, we create new functions or dataclasses that encapsulate the updated behavior.

Environment Updates

Instead of modifying a dictionary, we define a new function that returns the new value for the updated identifier and delegates to the original environment for all other identifiers:

```python
def env_extend(env: Environment, name: str, value: DVal) -> Environment:
    """Create new environment with an added binding"""
    def new_env(n: str) -> DVal:
        if n == name:
            return value
        return env(n)
    return new_env
```

This function returns a new environment that: - Returns `value` when asked for `name` - Delegates to the original environment for all other identifiers

This approach: - Preserves referential transparency - Enables easy implementation of lexical scoping - Facilitates reasoning about program behavior - Models the mathematical concept of function extension

State Updates

Similarly, state updates create new State objects with updated store functions:

```python
def state_update(state: State, location: Location, value: MVal) -> State:
    """Create new state with an updated value at given location"""
    def new_store(loc: Location) -> MVal:
        if loc == location:
            return value
        return state.store(loc)
    return State(store=new_store, next_loc=state.next_loc)
```

This function creates a new State object that: - Contains a new store function that returns the new value when asked for the specified location - Delegates to the original state's store function for all other locations - Preserves the next_loc value from the original state

Empty Environment and State

The primitives for creating empty environments and states define initial values:

```python
def empty_environment() -> Environment:
    """Create an empty environment function"""
    def env(name: str) -> DVal:
        raise ValueError(f"Undefined identifier: {name}")
    return env


def empty_memory() -> State:
    """Create an empty memory state"""
    def store(location: Location) -> MVal:
        raise ValueError(f"Undefined memory location: {location}")
    return State(store=store, next_loc=0)
```

Note that empty_memory returns a State dataclass initialized with an empty store function and next_loc set to 0.

Memory Allocation

In a complete interpreter, we use the State dataclass to implement memory allocation elegantly:

```python
def allocate(state: State, value: MVal) -> tuple[State, Location]:
    """Allocate a new memory location and store a value there.
    Returns the updated state and the new location."""
    location = state.next_loc
    new_state = state_update(state, location, value)
```

```
    # Return state with incremented next_loc and the allocated location
    return State(store=new_state.store, next_loc=location + 1), location
```

This function: 1. Gets the next available location from the state 2. Updates the store function to map this location to the provided value 3. Returns a new state with the incremented next_loc and the allocated location

By bundling the store function with the next_loc counter in our State dataclass, we maintain a purely functional approach while elegantly handling the allocation challenge. This design demonstrates how functional programming can manage state without side effects by making state changes explicit in the return values of functions.

Initial Environment Setup

In our functional implementation, the initial environment is built by starting with an empty environment and extending it with each operator:

```python
def create_initial_env() -> Environment:
    """Create an environment populated with standard operators"""
    env = empty_environment()
    env = env_extend(env, "+", add)
    env = env_extend(env, "-", subtract)
    env = env_extend(env, "*", multiply)
    env = env_extend(env, "/", divide)
    env = env_extend(env, "%", modulo)
    return env
```

This builds up the environment incrementally, adding each binding through functional extension.

Section 5: Environment-Based Interpretation

Traditional interpreters often use pattern matching on operators directly in the evaluation function. An environment-based approach takes a more abstract view, treating operators as first-class values in the environment.

Traditional Approach (from Chapter 3)

```python
def evaluate(ast: Expression) -> int:
    match ast:
        case Number(value):
            return value
        case BinaryExpression(op, left, right):
            left_value = evaluate(left)
            right_value = evaluate(right)
            match op:
                case "+":
                    return left_value + right_value
                case "-":
                    return left_value - right_value
                # ... other operations
```

Environment-Based Approach

```python
def evaluate(ast: Expression, env: Environment) -> MVal:
    match ast:
        case Number(value):
            return value
        case BinaryExpression(op, left, right):
```

```python
            try:
                # Get operator from environment
                operator = env_lookup(env, op)

                # Ensure it's a DenOperator
                if not isinstance(operator, Callable):
                    raise ValueError(f"{op} is not a function")

                # Evaluate operands and apply operator
                left_value = evaluate(left, env)
                right_value = evaluate(right, env)

                # Apply the operator to the evaluated operands
                return operator(left_value, right_value)
            except ValueError as e:
                raise ValueError(f"Evaluation error: {e}")
```

Benefits of the Environment-Based Approach

- **Extensibility**: New operators can be added to the environment without modifying the evaluator
- **First-class operations**: Operators are values that can be passed, returned, and manipulated
- **Consistent treatment**: Operators and other identifiers are handled uniformly
- **Semantic clarity**: The environment explicitly represents the mapping from names to meanings

Section 6: Implementing an Environment-Based Interpreter

Our `domains.py` file implements a complete environment-based interpreter:

1. **Define semantic domains**: Denotable and memorizable values
2. **Implement operators**: Define functions for arithmetic operations
3. **Create the environment**: Populate with standard operators
4. **Evaluate expressions**: Using the environment to look up operators
5. **REPL**: Interactive environment for testing the interpreter

Section 7: Extending the Interpreter

This approach makes it easy to extend the language with new features:

Adding New Operators

To add a new operator, simply define its function and add it to the environment:

```python
def power(x: int, y: int) -> int:
    return x ** y

# Extend environment
env = env_extend(create_initial_env(), "**", power)

### Adding Variables

To support variables, extend the AST with a variable node and update the evaluator:

```python
@dataclass
class Variable:
    name: str
```

```python
# Update Expression type
type Expression = Number | BinaryExpression | Variable
```

Update evaluate function

```python
def evaluate(ast: Expression, env: Environment) -> MVal:
    match ast:
        case Variable(name):
            value = env_lookup(env, name)
            # Additional check might be needed if variables can only be Numbers
            if not isinstance(value, int):
                raise ValueError(f"{name} is not a number")
            return value
        # ... existing cases
```

Additional Resources

- Denotational Semantics (Wikipedia)
- Programming Language Semantics (Stanford)
- Functional Programming and Lambda Calculus
- Environment and Store in Programming Languages

Section 8: Primitives for Environment and Memory

In the formal semantics of programming languages, we work with primitives that define how environments and memory operate. These primitives capture the essential operations needed to model variable bindings and memory allocation.

Locations as a Semantic Domain

Before discussing memory operations, we must recognize **Locations** as a fundamental semantic domain:

```python
type Location = int  # In our implementation, locations are integers
```

Locations (sometimes called addresses) are abstract entities that serve as references to memory cells. They are a semantic domain distinct from integers used in calculation, even though we might represent them as integers in an implementation. In a language's formal semantics, locations are opaque values that only make sense in the context of memory operations.

Memory (State) Primitives

Memory, also called the store or state, is represented by the State dataclass which contains both a store function mapping locations to memorizable values and a next_loc field. The core operations include:

1. **empty_memory**: Creates an initial, empty memory state

   ```python
   def empty_memory() -> State:
       """Create an empty memory state"""
       def store(location: Location) -> MVal:
           raise ValueError(f"Undefined memory location: {location}")
       return State(store=store, next_loc=0)
   ```

2. **update**: Modifies the memory at a specific location

   ```python
   def state_update(state: State, location: Location, value: MVal) -> State:
       """Create new state with an updated value at given location"""
   ```

```python
def new_store(loc: Location) -> MVal:
    if loc == location:
        return value
    return state.store(loc)
return State(store=new_store, next_loc=state.next_loc)
```

3. **lookup**: Retrieves a value from a location

```python
def state_lookup(state: State, location: Location) -> MVal:
    """Look up a value at a given location"""
    try:
        return state.store(location)
    except ValueError:
        raise ValueError(f"Undefined memory location: {location}")
```

Memory in programming languages has a maximum size, determined by hardware or system configuration. When a program needs more memory than available:

- In simple interpreters: An "out of memory" error occurs
- In modern operating systems: Memory expansion mechanisms activate, such as:
    - Virtual memory paging to disk
    - Heap expansion
    - Garbage collection (freeing unused memory)

Environment Primitives

The environment maps identifiers to denotable values. Its essential operations include:

1. **empty_environment**: Creates an initial, empty environment

```python
def empty_environment() -> Environment:
    """Create an empty environment function"""
    def env(name: str) -> DVal:
        raise ValueError(f"Undefined identifier: {name}")
    return env
```

2. **bind** (or **extend**): Adds a new binding to an environment

```python
def env_extend(env: Environment, name: str, value: DVal) -> Environment:
    """Create new environment with an added binding"""
    def new_env(n: str) -> DVal:
        if n == name:
            return value
        return env(n)
    return new_env
```

3. **lookup**: Retrieves a value bound to an identifier

```python
def env_lookup(env: Environment, name: str) -> DVal:
    """Look up an identifier in the environment"""
    try:
        return env(name)
    except ValueError:
        raise ValueError(f"Undefined identifier: {name}")
```

Memory and Environment in Language Semantics

The interaction between environments and memory is central to understanding language features:

- **Variables**: Bind identifiers to locations (in the environment) which then hold values (in memory)

- **Assignment**: Changes memory but not the environment (the variable still refers to the same location)
- **Scoping**: Creates nested environments with different bindings for the same identifier
- **Parameter passing**: Creates bindings between formal parameters and actual arguments

This conceptual separation allows language designers to reason clearly about the effects of operations and ensure language features interact correctly.

Section 9: Practical Implementation and Testing

Our environment-based interpreter includes practical components for execution and testing, demonstrating how theoretical concepts translate to code.

Parsing and AST Construction

In our implementation, we use the Lark parser to convert text expressions into parse trees:

```python
def parse_ast(expression: str) -> Expression:
    """Parse a string expression into an AST"""
    parse_tree = parser.parse(expression)
    return transform_parse_tree(parse_tree)
```

The `transform_parse_tree` function then converts these parse trees into our AST representation, ready for evaluation.

Interactive REPL

A Read-Eval-Print Loop allows interactive testing of the interpreter:

```python
def REPL():
    """Read-Evaluate-Print Loop with environment"""
    env = create_initial_env()

    print("Mini-interpreter with environment (type 'exit' to quit)")
    print("Available operators: +, -, *, /, %")

    while True:
        expression = input("Enter an expression (exit to quit): ")
        if expression == "exit":
            break

        try:
            ast = parse_ast(expression)
            result = evaluate(ast, env)
            print(result)
        except Exception as e:
            print(f"Error: {e}")
```

The REPL shows how our environment-based interpreter integrates with user interaction.

Automated Tests

Testing ensures the interpreter behaves as expected across various expressions:

```python
def run_tests():
    """Run test expressions to verify the parser and evaluator"""
    test_expressions = [
        "1+2",
        "3*4",
        "5-3",
```

```
        "10/2",
        "10%3",
        "(1+2)*3",
        "1+(2*3)",
        "10/(2+3)",
        "10%(2+3)",
    ]

    env = create_initial_env()

    for expr in test_expressions:
        try:
            ast = parse_ast(expr)
            result = evaluate(ast, env)
            print(f"{expr} = {result}")
        except Exception as e:
            print(f"{expr} -> Error: {e}")
```

These tests validate core functionality while providing examples of valid expressions.

Section 10: Conclusion and Next Steps

Summary of Key Concepts

In this chapter, we've explored:

1. **Semantic domains** as mathematical structures that give meaning to programs
2. **Environment-based interpretation** as a flexible approach to language implementation
3. **Functional representations** of environments and state
4. **Primitive operations** for manipulating environments and memory
5. **The distinction** between denotable and memorizable values

These concepts form the foundation for understanding more complex language features in subsequent chapters.

Looking Forward

The environment-based approach introduced here will be extended in future chapters to support:

1. **Variables and assignment**: Using environments to bind identifiers to memory locations
2. **Scoping mechanisms**: Creating nested environments for block-structured code
3. **Functions and closures**: Capturing environments for later execution
4. **Typing systems**: Adding constraints on what values expressions can produce

By building on these semantic foundations, we can construct a rich understanding of programming language design and implementation.

Exercises for the Reader

1. Extend the interpreter to support variables using the Variable AST node described in Section 7
2. Add support for multi-character operators (e.g., "**" for exponentiation)
3. Implement a memory system with the primitives described in Section 8
4. Add support for conditional expressions (if-then-else)
5. Implement a simple block structure with local variables

These exercises will deepen your understanding of language semantics and interpreter design while preparing you for more advanced topics in the next chapters.